

Adversarial Attacks

Robust Deep Learning (Seminar)

Bjarne Thal Marius Wrobel

Dept. of Computer Science, RWTH Aachen University, Germany
`{bjarne.thal, marius.wrobel}@rwth-aachen.de`

August 4, 2026



Which one is the gibbon?



Which one is the gibbon?



panda



gibbon

Adversarial Examples



“panda”
57.7% confidence

+ .007 ×




=



“gibbon”
99.3 % confidence




Goodfellow et al., 2014

Adversarial Examples



Transferability




- ▶ Attacks transfer across architectures, training data, and more



Papernot et al., 2016

Transferability

- ▶ Attacks transfer across architectures, training data, and more
- Adv. ex. are properties of a problem itself



Papernot et al., 2016

Why this matters

- ▶ **Fundamental blind spot** while integrated AI systems require high reliability



Images: medicalimaging.na; thetravelimages.com

Outline

1. Evasion attacks
2. Poisoning attacks
3. Discussion
 - 3.1 Explaining adversarial examples
 - 3.2 Adversarial attacks as a subject of study

Outline

1. Evasion attacks
2. Poisoning attacks
3. Discussion

Taxonomy & Terminology

Attacker knowledge

White-box

vs.

Black-box

Attack stage

Evasion (at test time)

vs.

Poisoning (at training time)

Attack goal

Targeted (specific wrong label)

vs.

Untargeted (any wrong label)

L-BFGS

$$\min_{\delta} \|\delta\|_2 \quad \text{subject to} \quad f(x + \delta) = t$$

- ▶ f has no gradient — it returns a discrete class label, $f(x + \delta) = t$ is a hard, non-differentiable condition
 - ▶ Derivative-based optimisers such as L-BFGS need a gradient to move at all
- ↪ Reformulate

$$\min_{\delta} c \|\delta\|_2 + L_f(x + \delta, t)$$

Significance

- ▶ First method to **reliably generate** adversarial examples
- ▶ Computationally expensive — far too slow to be practical
- ↪ It *proved a point* rather than providing a usable attack

Fast Gradient Sign Method (FGSM)

Fast Gradient Sign Method

Goodfellow et al., 2014

White-box · Evasion · Untargeted

- ▶ Previously: Adversarial examples assumed to be caused by 'complex' behaviour
- ▶ Goodfellow et al.: Linear models have adv. ex. as well
- Models may be too linear

Do Neural Networks behave linearly?

Goodfellow et al., 2014

...

Fast Gradient Sign Method

Goodfellow et al., 2014

White-box · Evasion · Untargeted

- ▶ Linearise the loss function:

$$\arg \max_{\|\delta\|_{\infty} \leq \epsilon} L_f(x + \delta, y) = \arg \max_{\|\delta\|_{\infty} \leq \epsilon} L_f(x, y) + \nabla_x L_f(x, y)^T \delta$$

- ▶ Maximised by $\delta = \epsilon \cdot \text{sign}(\nabla_x L_f(x, y))$

↪ Single step in gradient direction

Fast Gradient Sign Method - Example

Goodfellow et al., 2014

White-box · Evasion · Untargeted

placeholder: step graph here

[Results/Significance]

Goodfellow et al., 2014

White-box · Evasion · Untargeted


- ▶ Maximised by $\delta = \epsilon \cdot \text{sign}(\nabla_x L_f(x, y))$

↪ Single step in gradient direction

High-dimensional linearity

Goodfellow et al., 2014

- Affine case first: perturbation shifts activation by $w^\top \eta$



$$\begin{aligned}w^\top \tilde{x} &= w^\top (x + \eta) \\ &= w^\top x + w^\top \eta\end{aligned}$$


$$w^\top \eta = \epsilon \sum_{i=1}^n |w_i| = \epsilon m n$$

n dimensions, mean weight magnitude m

High-dimensional linearity

Goodfellow et al., 2014

- ▶ Affine case first: perturbation shifts activation by $w^\top \eta$
- ▶ Best strategy: η aligned with the weights



$$\begin{aligned}w^\top \tilde{x} &= w^\top (x + \eta) \\&= w^\top x + w^\top \eta\end{aligned}$$


$$w^\top \eta = \epsilon \sum_{i=1}^n |w_i| = \epsilon m n$$

n dimensions, mean weight magnitude m

High-dimensional linearity

Goodfellow et al., 2014

- ▶ Affine case first: perturbation shifts activation by $w^\top \eta$
- ▶ Best strategy: η aligned with the weights
- ▶ Score change scales *linearly* with number of dimensions w.r.t. input change



$$\begin{aligned}w^\top \tilde{x} &= w^\top (x + \eta) \\&= w^\top x + w^\top \eta\end{aligned}$$


$$w^\top \eta = \epsilon \sum_{i=1}^n |w_i| = \epsilon m n$$

n dimensions, mean weight magnitude m

High-dimensional linearity

Goodfellow et al., 2014

- ▶ Affine case first: perturbation shifts activation by $w^\top \eta$
 - ▶ Best strategy: η aligned with the weights
 - ▶ Score change scales *linearly* with number of dimensions w.r.t. input change
- ~> Many imperceptible changes accumulate into one large shift



$$\begin{aligned}w^\top \tilde{x} &= w^\top (x + \eta) \\ &= w^\top x + w^\top \eta\end{aligned}$$

$$w^\top \eta = \epsilon \sum_{i=1}^n |w_i| = \epsilon m n$$

n dimensions, mean weight magnitude m

From linear models to deep networks

Goodfellow et al., 2014

- ▶ Most DNN architectures behave *locally linearly* for ease of optimisation

ReLU

LSTM

sigmoid near zero

maxout

linear-softmax

From linear models to deep networks

Goodfellow et al., 2014

- ▶ Most DNN architectures behave *locally linearly* for ease of optimisation
- ↪ Within ϵ -ball, the linear argument carries over

ReLU

sigmoid near zero

LSTM

maxout

linear-softmax

Projected Gradient Descent (PGD)

General formulation of adversarial training [...! incl. defenders part?]. Adversary solves

$$\max_{\delta \in S} L_f(x + \delta, y)$$

- ▶ FGSM: Use $S = \{\delta \mid \|\delta\|_{\infty} \leq \epsilon\}$
- ▶ Use projected gradient descent to optimise

↪ Models may be too linear

PGD - a general first-order attack?

Madry et al., 2017

White-box · Evasion · ???


DeepFool

DeepFool

Moosavi-Dezfooli et al., 2016 - the affine, binary case

White-box · Evasion · Untargeted

- ▶ Decision boundary is a *hyperplane*
- ▶ Minimal perturbation is *orthogonal projection* of x onto $\mathcal{F} \rightarrow$ closed form




DeepFool

Moosavi-Dezfooli et al., 2016 - the general, multiclass case

White-box · Evasion · Untargeted

- ▶ **Non-linear:** iterative $\rightarrow$ linearise the score function at current point (first-order Taylor), project onto that hyperplane, repeat until label flips
- ▶ **Multiclass:** Linearise every class boundary, project onto the **closest**



Significance

- ▶ **Geometric reading:** robustness as distance to the decision boundary

Transfer attacks

Transfer attacks

- ▶ Black-box attack pattern
- ▶ Adversary creates and trains a substitute model for the target
- ▶ Adversarial examples for substitute model are found (white-box!)
- ▶ These can be used to attack the target

Jacobian-Based Dataset Augmentation

Papernot et al., 2017

Black-box · Evasion · ???

- ▶ Assumption: No knowledge of model parameters, dataset
Access to somewhat representative inputs is granted (e.g. natural images)
- ▶ How do we construct a suitable substitute?
- ▶ Construct an artificial dataset!

Jacobian-Based Dataset Augmentation

Papernot et al., 2017

Black-box · Evasion · ???

- ▶ Assumption: No knowledge of model parameters, dataset
Access to somewhat representative inputs is granted (e.g. natural images)
- ▶ How do we construct a suitable substitute? (similar decision boundaries)
- ▶ Construct an artificial dataset!

Jacobian-Based Dataset Augmentation

Papernot et al., 2017

Black-box · Evasion · ???

- ▶ Use target model as an oracle to label samples
- ▶ Strategy for sample selection is necessary
- ▶ Idea: iteratively add meaningful samples to dataset
- ▶ Use Jacobian for a heuristic

Heuristic for meaningful directions

Papernot et al., 2017

Black-box · Evasion · ???

$$x = x + \alpha \cdot \text{sign}(J_f(x)[o(x)])$$

- ▶ $J_f(x)$: total derivative
- ▶ $J_f(x)[o(x)]$: total derivative, column of the correct label
- We move into the class region (according to the substitute)!

Jacobian-Based Dataset Augmentation

Papernot et al., 2017

Black-box · Evasion · ???

- ▶ 1) Collect samples somewhat related to the input domain
- ▶ 2) Choose architecture for substitute
- ▶ 3) Target model labels current samples
- ▶ 4) Train substitute
- ▶ 5) Expand dataset: $S_{N+1} := \{x = x + \alpha \cdot \text{sign}(J_f(x)[o(x)]) \mid x \in S_N\} \cup S_N$
- ▶ Repeat (3) onwards

White-box attacks to improve transferability

- ▶ Adversarial examples can differ in transferability!
- ▶ MI-FGSM: adds momentum to FGSM
- ▶ Ensemble attacks

ZOO

ZOO: Zeroth order black-box attacks

Black-box · Evasion · Targ./Untarg.

- ▶ Black-box attack without substitute model
- ▶ No quality loss due to transfer
- ▶ Assumptions: Adversary allowed to query, gets confidence scores

$$\hat{g}_i := \frac{\partial l(x)}{\partial x_i} \approx \frac{l(x + he_i) - l(x - he_i)}{2h}$$

- ▶ l loss function based on confidence scores
- ▶ Two queries per coordinate, full gradient infeasible
- ▶ -! Coordinate descent using first and/or second-order method

- ▶ Reminder: Carlini-Wagner loss: $\ell(x') = \max(\max_{i \neq t} z_i(x') - z_t(x'), 0)$
- ▶ Targeted ZOO: $l(x, t) = \max\{\max_{i \neq t} \{\log[f(x)]_i\} - \log[f(x)]_t, -\kappa\}$
- ▶ Untargeted ZOO: $l(x) = \max\{\log[f(x)]_y - \max_{i \neq y} \{\log[f(x)]_i\}, -\kappa\}$
- ▶ *log* to estimate logits

Outline

1. Evasion attacks
2. Poisoning attacks
3. Discussion

Hidden Trigger Backdoor Attack

Hidden Trigger Backdoor Attack

Black-box · ??? ·

- ▶ Goal: get ML model to misclassify any sample with trigger pattern into target class t
- ▶ Similar threat model to Poison Frogs
- ▶ Restriction: adversary cannot choose labels
- ▶ -i "Clean-label attack"

- ▶ Sample instances s b of target class t
- ▶ Apply trigger to source: patched source $\tilde{s}$
- ▶ Poison instances:

$$\arg \min_x ||g(x) - g(\tilde{s})||_2^2 \quad \textbf{s.t.} \quad ||x - b||_\infty < \epsilon$$

- ▶ g : feature representation (second-to-last layer output)
- ▶ Keep x close to $\tilde{s}$ in feature space, but close to b in pixel space

Outline

1. Evasion attacks
2. Poisoning attacks
3. Discussion
 - 3.1 Explaining adversarial examples
 - 3.2 Adversarial attacks as a subject of study

Explaining adversarial examples

2013 Szegedy et al. — quirks of non-linearity

Explaining adversarial examples

2013 Szegedy et al. — quirks of non-linearity

2014 Goodfellow et al. — high-dimensional linearity

Explaining adversarial examples

2013 **Szegedy et al.** — quirks of non-linearity

2014 **Goodfellow et al.** — high-dimensional linearity

2019 **Ilyas et al.** — predictive but non-robust features

Predictive and robust features

Ilyas et al., 2019

Predictive

Correlated with the label in expectation.

Predictive and robust features

Ilyas et al., 2019

Predictive

Correlated with the label in expectation.

Robust

Stays correlated *under* adversarial perturbation.

Predictive and robust features

Ilyas et al., 2019

Predictive

Correlated with the label in expectation.

Robust

Stays correlated *under* adversarial perturbation.

→ **Predictive but non-robust:** Standard training happily relies on such features

Experiment: "Misabeled" training dataset

Ilyas et al., 2019



Real-life adversarial attacks

- ▶ Real confirmed malicious attacks are surprisingly rare
- ▶ Two candidate explanations:
 1. They slip through unnoticed, logged as ordinary misclassifications
 2. They are simply of little use to actual attackers

Adversarial attacks as a subject of study

- ▶ Not really significant security threat

Adversarial attacks as a subject of study

- ▶ Not really significant security threat
- ▶ Mainly *epistemic significance*:
 - ▶ Not: What can an attacker do?
 - ▶ But: What do they reveal about the targetted system?

Adversarial attacks as a subject of study

- ▶ Not really significant security threat
- ▶ Mainly *epistemic significance*:
 - ▶ Not: What can an attacker do?
 - ▶ But: What do they reveal about the targetted system?

Before

- ▶ Accuracy as near-single metric
- ▶ Robustness assumed implicitly

Adversarial attacks as a subject of study

- ▶ Not really significant security threat
- ▶ Mainly *epistemic significance*:
 - ▶ Not: What can an attacker do?
 - ▶ But: What do they reveal about the targetted system?

Before

- ▶ Accuracy as near-single metric
- ▶ Robustness assumed implicitly

Afterwards

- ▶ High accuracy $\neq$ human-like understanding
- ▶ Robustness as property of interest

Key takeaways

- ▶ Imperceptible perturbation → wrong prediction
- ▶ Transferable across models, architectures, training data
- ▶ Categories: white/black-box · evasion/poisoning · targeted/untargeted
- ▶ Phenomenon still not fully explained
- ▶ Marginal security threat so far but epistemic relevance

Thank you for your attention!

Questions?



References I

Backup Slides

Why that works

Ilyas et al., 2019

	Robust features point at	Non-robust features point at	Label	Image
Original clean training image (in $\mathcal{D}$)	dog	dog	dog	
Perturbed training image (in $\hat{\mathcal{D}}_{NR}$)	dog <i>(unaffected by perturbation)</i>	cat <i>(through adversarial perturbation)</i>	cat	
Original clean test image (in $\mathcal{D}$)	cat	cat	cat (predicted)	


What is an Adversarial Example?

Definition (Adversarial Example)

An *adversarial example* is an **input to a machine learning model** that has been intentionally perturbed, typically **in a subtle and often imperceptible manner**, such that the model **produces an incorrect prediction** while the semantic content of the input remains unchanged.

- ▶ **Machine learning models:** No quirk of deep networks, but a general property of machine learners
- ▶ **Subtle and imperceptible perturbations:** mathematically small under a chosen norm / visually imperceptible to humans
- ▶ **Produces an incorrect prediction:** The machine learner is fooled where the human would not have been

Measuring “small”: the ℓ_p norms



Norm defines the **perturbation budget** ϵ .

$$\underbrace{\min_{\delta} c \|\delta\|_2 + L_f(x + \delta, t)}_{\text{L-BFGS}} \longrightarrow \underbrace{\min_{\delta} \|\delta\|_p + c \cdot \ell(x + \delta)}_{\text{C\&W}}$$

- ▶ $\ell(x') = \max(\max_{i \neq t} z_i(x') - z_t(x'), 0)$
- ▶ Built so that $\ell(x') \leq 0 \iff$ input is classified as $t \rightarrow$ faithfully encodes the attack goal
- ▶ c decided through **binary search**

$$x' = \frac{1}{2}(\tanh(w) + 1), \quad w \in \mathbb{R}^m$$

- ▶ Mapping any $w \in \mathbb{R}^m$ into $(0, 1)^m \rightarrow$ every w decodes to a **valid image** $\rightarrow$ dissolving box-constraint
- ▶ Optimise w **unconstrained** with Adam optimiser

Significance

- ▶ Strongest attack of its time across $\ell_0/\ell_2/\ell_\infty$
- ▶ Broke *defensive distillation*, exposing it as *gradient masking*

Physical adversarial attacks



Sharif et al., 2016 · Xu et al., 2020



Eykholt et al., 2018



Athalye et al., 2018

- ▶ Goal: get ML model to misclassify one specific sample v into target class t
- ▶ Assumption: model based on pretrained weights known to the attacker, access to data distribution
- ▶ Restriction: adversary cannot choose labels
- ▶ -! "Clean-label attack"

- ▶ Sample instance b of target class t
- ▶ Poison instance(s):

$$p = \arg \min_x \|g(x) - g(v)\|_2^2 + \lambda \|x - b\|_2^2$$

- ▶ g : feature representation (second-to-last layer output)
- ▶ Keep x close to v in feature space, but close to b in pixel space