

Adversarial Attacks

Robust Deep Learning (Seminar)

Bjarne Thal Marius Wrobel

Dept. of Computer Science, RWTH Aachen University, Germany
`{bjarne.thal, marius.wrobel}@rwth-aachen.de`

August 4, 2026



Two images



Two images



panda



gibbon

A fragile foundation



x

“panda”

57.7% confidence

$+ .007 \times$



$\text{sign}(\nabla_x J(\theta, x, y))$

“nematode”

8.2% confidence

$=$



$x +$

$\epsilon \text{sign}(\nabla_x J(\theta, x, y))$





“gibbon”

99.3 % confidence

Goodfellow et al., 2014





Transferability

- An adversarial example crafted against *one* model often fools models it was never shown.







Transferability

- ▶ An adversarial example crafted against *one* model often fools models it was never shown.
- ▶ This survives a change of **architecture** — and of **training data**, initialisation, and hyperparameters.







Transferability

- ▶ An adversarial example crafted against *one* model often fools models it was never shown.
- ▶ This survives a change of **architecture** — and of **training data**, initialisation, and hyperparameters.
- ▶ This points to something deeper: the vulnerability appears to be **inherent to the problem being learned**, not a defect of a particular model, or even of a particular architecture.



Transferability

- ▶ An adversarial example crafted against *one* model often fools models it was never shown.
 - ▶ This survives a change of **architecture** — and of **training data**, initialisation, and hyperparameters.
 - ▶ This points to something deeper: the vulnerability appears to be **inherent to the problem being learned**, not a defect of a particular model, or even of a particular architecture.
- ↪ The attacker needs **no access** to the target. Train your own model, attack that, and the perturbation carries over.



Why this matters

- ▶ What we just saw exposes a **fundamental blind spot** — the model and the human are not looking at the same thing.



Images: medicalimaging.na; evsmarts.com; thetravelimages.com

Why this matters

- ▶ What we just saw exposes a **fundamental blind spot** — the model and the human are not looking at the same thing.
- ▶ That raises hard questions about the **reliability** of systems we already treat as dependable.



Images: medicalimaging.na; evsmarts.com; thetravelimages.com

Why this matters

- ▶ What we just saw exposes a **fundamental blind spot** — the model and the human are not looking at the same thing.
- ▶ That raises hard questions about the **reliability** of systems we already treat as dependable.
- ▶ And it grows more pressing as these systems are woven ever deeper into **everyday life**, where failures carry real consequences.



Images: medicalimaging.na; evsmarts.com; thetravelimages.com

Outline

1. What is an adversarial example?
2. Taxonomy
3. Evasion attacks
4. Poisoning attacks
5. Discussion
 - 5.1 Physical adversarial attacks
 - 5.2 Real-life adversarial attacks
 - 5.3 Explaining adversarial examples
 - 5.4 Adversarial attacks as a subject of study

Outline

1. What is an adversarial example?
2. Taxonomy
3. Evasion attacks
4. Poisoning attacks
5. Discussion

Definition

Definition 1 (Adversarial Example)

An *adversarial example* is an **input to a machine learning model** that has been **intentionally perturbed**, typically **in a subtle and often imperceptible manner**, such that the model produces an incorrect prediction while the semantic content of the input remains unchanged.


Definition

Definition 1 (Adversarial Example)


An *adversarial example* is an **input to a machine learning model** that has been **intentionally perturbed**, typically **in a subtle and often imperceptible manner**, such that the model produces an incorrect prediction while the semantic content of the input remains unchanged.

- ▶ **Machine learning models** — no quirk of deep networks, but a general property of machine learners.
- ▶ **Intentional perturbations** — generated systematically through optimisation and exploited by an adversary, not stumbled upon.
- ▶ **Subtle and imperceptible perturbations** — mathematically small under a chosen norm, and visually indistinguishable to humans; the latter is asserted far more often than it is rigorously validated.


Measuring “small”: the ℓ_p norms



$$\|x\|_0 = (|x_1|^0 + \dots + |x_n|^0)$$



$$\|x\|_2 = (|x_1|^2 + \dots + |x_n|^2)^{\frac{1}{2}}$$



$$\|x\|_\infty = \max_i |x_i|$$

The norm defines the **perturbation budget** ϵ and, with it, what “similar” even means.

Outline

1. What is an adversarial example?
2. Taxonomy
3. Evasion attacks
4. Poisoning attacks
5. Discussion

Four orthogonal axes

- ▶ **Attacker knowledge** — white-box vs. black-box
- ▶ **Attack stage** — evasion (test time) vs. poisoning (train time)
- ▶ **Attack goal** — targeted vs. untargeted
- ▶ **Perturbation constraint** — which norm, which budget

Outline

1. What is an adversarial example?
2. Taxonomy
3. Evasion attacks
4. Poisoning attacks
5. Discussion

$$\min_{\delta} \|\delta\|_2 \quad \text{subject to} \quad f(x + \delta) = t$$

- ▶ f has no gradient — it returns a discrete class label, $f(x + \delta) = t$ is a hard, non-differentiable condition.
 - ▶ Derivative-based optimisers such as L-BFGS need a gradient to move at all.
- ↪ Reformulate.

$$\min_{\delta} c \|\delta\|_2 + L_f(x + \delta, t)$$

Significance


- ▶ First method to **reliably generate** adversarial examples.
- ▶ Computationally expensive — far too slow to be practical.
- ↪ It *proved a point* rather than providing a usable attack.

The arc of the field

2013 **L-BFGS** — expensive, but proves existence

- ▶ For an affine $f(x) = w^\top x + b$, the decision boundary $\mathcal{F}$ is a **hyperplane**.
- ▶ Minimal perturbation is the **orthogonal projection** of x_0 onto $\mathcal{F}$ — closed form, no search.

$$r_*(x_0) = -\frac{f(x_0)}{\|w\|_2^2} w$$




DeepFool

The general, multiclass case

White-box · Evasion · Untargeted · ℓ_2

- ▶ **Non-linear:** iterative - linearise the score function at current point (first-order Taylor), project onto that hyperplane, repeat until label flips.
- ▶ **Multiclass:** Linearise every class boundary, project onto the **closest**.



Significance

- ▶ Gave problem a **geometric reading**: robustness as distance to the decision boundary.
- ▶ First to **extend to any ℓ_p norm**, $p \in [1, \infty)$.

The arc of the field

2013 **L-BFGS** — expensive, but proves existence

2016 **DeepFool** — minimal perturbation via linear approximation

$$\underbrace{\min_{\delta} c \|\delta\|_2 + L_f(x + \delta, t)}_{\text{L-BFGS}} \longrightarrow \underbrace{\min_{\delta} \|\delta\|_p + c \cdot \ell(x + \delta)}_{\text{C\&W}}$$

- ▶ $\ell(x') = \max(\max_{i \neq t} z_i(x') - z_t(x'), 0)$.
- ▶ Built so that $\ell(x') \leq 0 \iff$ input is classified as t - ℓ faithfully encodes the attack goal.
- ▶ c decided through **binary search**

$$x' = \frac{1}{2}(\tanh(w) + 1), \quad w \in \mathbb{R}^m$$

- ▶ mapping any $w \in \mathbb{R}^m$ into $(0, 1)^m$ - $\hookrightarrow$ every w decodes to a **valid image** - $\hookrightarrow$ dissolving box-constraint
- ▶ Optimise w **unconstrained** with Adam optimiser

Significance

- ▶ Strongest attack of its time across $\ell_0/\ell_2/\ell_\infty$.
- ▶ Broke *defensive distillation*, exposing it as *gradient masking*

The arc of the field

2013 **L-BFGS** — expensive, but proves existence

2016 **DeepFool** — minimal perturbation via linear approximation

2017 **C&W** — reformulate as optimisation; breaks defences

Outline

1. What is an adversarial example?
2. Taxonomy
3. Evasion attacks
4. Poisoning attacks
5. Discussion

Outline

1. What is an adversarial example?
2. Taxonomy
3. Evasion attacks
4. Poisoning attacks
5. Discussion
 - 5.1 Physical adversarial attacks
 - 5.2 Real-life adversarial attacks
 - 5.3 Explaining adversarial examples
 - 5.4 Adversarial attacks as a subject of study

Physical adversarial attacks

Real-life adversarial attacks

Explaining adversarial examples

Adversarial attacks as a subject of study

Take-home messages

References I

Backup: attack comparison